


Clearly indicate the edges of the following spanning trees of the weighted graph by, say, drawing a new copy of the graph with vertex labels and edge weights intact.

Each part is worth 2.5 points out of 10.

(a) A depth-first spanning tree rooted at s


Solution:



■

(b) A breadth-first spanning tree rooted at s


Solution:



■

(c) A shortest paths tree rooted at s


Solution:



■

(d) A minimum spanning tree

Solution:



■

Suppose you are a shopkeeper living in a country with n different types of coins, with values $1 = c[1] < c[2] < \dots < c[n]$. Not wanting to burden their pockets, whenever you give a customer change, you always use the smallest possible number of coins.

- (a) **(3 out of 10)** Show there is a set of coin values for which the greedy algorithm does *not* always give the smallest possible number of coins. Please state both the coin values and the amount of change C for which the greedy algorithm is not optimal.

Solution: Consider the coin values $\{1, 5, 6\}$. If asked to make $C = 10$ units of the changes, the greedy algorithm would use five coins: $\langle 6, 1, 1, 1, 1 \rangle$. However, it is possible to use only two coins: $\langle 5, 5 \rangle$. ■

- (b) **(7 out of 10)** Suppose your country's government decides to impose a currency system where the coin denominations are consecutive powers $b^0, b^1, b^2, \dots, b^k$ for some integer $b \geq 2$. Prove the greedy algorithm described in part (a) does make optimal change in this currency system.

Solution: First, observe that for any $0 \leq i \leq k - 1$, the value b^{i+1} is an integer multiple of b^i . Inductively, one cannot exceed the value b^i using smaller coins without meeting it, so one cannot exceed b^{i+1} using smaller coins without meeting it as well.

Now, suppose we're asked to make change C , and let i be the largest integer between 0 and k such that $b^i \leq C$. Suppose the optimal method does not use coin b^i . From the above, we see one can add up the coins in the optimal set in any order, and at a certain moment meet b^i exactly. We can replace all the coins up to this moment with the single coin b^i , contradicting the optimality of the coin set that excludes b^i .

Finally, we see the optimal change for C includes b^i . Inductively, the greedy algorithm goes on to make the remaining change $C - b^i$ optimally. ■

- (a) **(3 out of 10)** Let $G = (V, E)$ be a connected undirected graph with real edge weights $w : E \rightarrow \mathbb{R}$. Briefly describe and analyze a fast algorithm to compute a *maximum* weight spanning tree of G . You may assume edge weights are distinct. You do not need to justify correctness of your algorithm.

Solution: We define the edge weight function $w' : E \rightarrow \mathbb{R}$ where $w'(e) := -w(e)$. Then, we compute and return a *minimum* spanning tree under weight function w' using any algorithm from class such as Kruskal's. Computing the minimum spanning tree is the bottleneck, so the algorithm runs in $O(E \log V)$ time. ■

Explanation Any spanning tree T of weight $w(T)$ under w weights $-w(T)$ under w' . Therefore, minimizing the weight of T under w' is equivalent to maximizing it under w .

- (b) **(7 out of 10)** Let $G = (V, E)$ be a connected undirected graph with *positive* edge weights $w : E \rightarrow \mathbb{R}_{>0}$. A **feedback edge set** of G is a subset F of its edges such that every cycle in G contains at least one edge in F . In other words, removing every edge in F turns the graph G into a forest. Briefly describe and analyze a fast algorithm to compute a minimum weight feedback edge set of G . You may assume edge weights are distinct, and you may assume your solution to part (a) is correct. You do not need to justify correctness of your algorithm.

Solution: We compute a maximum weight spanning tree T using the algorithm from part (a) and then return $E - T$. The algorithm from part (a) is the bottleneck, so this procedure takes $O(E \log V)$ time as well. ■

Explanation Let F be a minimum weight feedback edge set. The graph $G - F$ is connected; otherwise, we could take any edge e connecting two components of $E \setminus F$ and remove it from F , reducing the weight of F while still leaving $G - F$ as a forest. Therefore, $G - F$ is a spanning tree. Minimizing the weight of feedback edge set F is equivalent to maximizing the weight of spanning tree $G - F$.

Suppose we are given a directed acyclic graph $G = (V, E)$ with real edge lengths $\ell : E \rightarrow \mathbb{R}$ such that G has a unique source s and a unique sink t . In other words, s is the only vertex with no predecessor and t is the only vertex with no successor. We have also been given a non-negative integer k and told that some vertices of G have been marked *important*. For simplicity, we may assume neither s nor t have been marked important.

Our goal is to design an algorithm to find the maximum length over all paths from s to t that contain at least k important vertices.

- (a) **(5 out of 10)** For any vertex $v \in V$ and any non-negative integer r , let $MaxLength(v, r)$ denote the maximum length over all paths from v to t that contain at least r important vertices. If no such path exists, let $MaxLength(v, r) := -\infty$.

Give a recursive definition of $MaxLength(v, r)$. You do not need to justify correctness of your definition.

Solution: For simplicity, we define a max over no elements to be $-\infty$.

$$MaxLength(v, r) = \begin{cases} 0 & \text{if } v = t \text{ and } r = 0 \\ -\infty & \text{if } v = t \text{ and } r > 0 \\ \max_{v \rightarrow w} (\ell(v \rightarrow w) + MaxLength(w, 0)) & \text{if } v \neq t \text{ and } r = 0 \\ \max_{v \rightarrow w} (\ell(v \rightarrow w) + MaxLength(w, r)) & \text{if } v \neq t, r > 0 \text{ and } v \text{ is not important} \\ \max_{v \rightarrow w} (\ell(v \rightarrow w) + MaxLength(w, r - 1)) & \text{otherwise} \end{cases}$$

■

- (b) **(5 out of 10)** Describe and analyze an efficient dynamic programming algorithm to compute the maximum length over all paths from s to t that contain at least k important vertices. If no such path exists, your algorithm should return $-\infty$.

Solution: We want to solve $MaxLength(v, r)$ for all $v \in V$ and values $0 \leq r \leq k$. Each subproblem $MaxLength(v, r)$ depends upon those taking a vertex w such that $v \rightarrow w$ is an edge. Therefore, we can iterate over all vertices in postorder and all r in arbitrary order. We return the value $MaxLength(s, k)$.

For each of the $O(k)$ choices of r , we solve one subproblem per vertex and have one dependency per edge. Therefore, the algorithm will run in $O(k(V + E))$ time. ■